

SUMMARY

- Senior backend engineer with 7+ years at consumer and fintech companies, now building LLM systems at Mercari, from a production Go/GCP SEO pipeline to prompt benchmarks and agent context and memory.
- 449 merged PRs across 29 services of a Go e-commerce monorepo (3rd most active of 87 committers in 2026), spanning seller and buyer features, platform migrations, and a partner-facing public API.
- Four internal Mercari talks on building MCP servers, harness engineering, Claude Code subagents vs. workflows, and the token cost of context compaction.

EDUCATION

- **East China Normal University** Shanghai, China
Bachelor of Software Engineering Jul 2015 – Jun 2019

EXPERIENCE

- **Mercari** Tokyo, Japan
Senior Software Engineer Aug 2024 – Present

AI & LLM Engineering

- **LLM SEO pipeline:** built a Go pipeline (gRPC, Pub/Sub workers, Gemini) for listing SEO copy; rolled it out to 100% after an analyst's 50/50 A/B test lifted impressions 93% and clicks 12% per listing.
- **Prompt engineering:** rewrote the SEO prompt as 7 grounded rules and a worked example, chosen on a self-built 334-call benchmark (invented facts on trap cases 5% to 0%, genuine quantities kept 98%).
- **Agent context benchmarking:** piloted an agent knowledge base (right doc found 5/5, cited 6/6) and wrote its 12 docs; separately cut 19k unused tool tokens per AI PR review and the skill list 32%.
- **Context management:** built a Claude Code agent on an internal dev machine that drafts partner API inquiry replies for human review in a median 7 min, with untrusted-thread fencing and 7 scoped tools.
- **Memory layer:** wired the agent to the claude-memory MCP server (see Projects); found 16 of 19 records never retrieved and moved to one query per key term, which matched 61 of 63 logged queries on replay.

Product Systems

- **Time sale:** led the backend redesign for campaigns of 4M+ discounted items, with 300-item batched starts every 15 min (2.5x stress-test throughput) and a unified bulk CSV pipeline built with 3 teammates.
- **Refurbished listings:** owned the backend (design doc, 10 delegated tasks) through a staged launch, generated the category set (+125 categories), and built the refurbished flag a new storefront reuses.
- **Storefront page migration:** led the backend move of seller storefront pages to the main marketplace domain for SEO: 2 design docs, 10 planned tasks (8 delegated), a review RPC, and 4 list RPCs.

Platform & Reliability

- **Valkey migration:** consolidated 21 over-provisioned Redis instances (60% of memory idle) into 6 shared HA Valkey clusters with IAM auth, cutting over one service at a time; audited net cache spend fell 53%.
- **Shared cache client:** rewrote the Go client on valkey-go (multiplexing, client-side caching) and moved 18 services onto it; on the 3 measured services, connections fell 29–59% and churn 65–77% vs. a v1 control.
- **Category service:** built an in-memory category-tree index reused by coupon and product services; cached a licensed-category check behind 87% of the service's SQL, cutting SQL per request nearly 80%.

- ROMS** Tokyo, Japan
Senior Software Engineer *Oct 2023 – Aug 2024*
 - Served as technical lead for a serverless, event-driven gateway platform connecting customer applications and warehouse hardware systems.
 - Defined system architecture and specifications supporting thousands of daily events.
 - Delivered MVP systems and internal tools, reducing time-to-demo from months to weeks.
- PayPay (Yahoo Group)** Tokyo, Japan
Backend Engineer *Jun 2023 – Sep 2023*
 - Improved accounting-platform reliability by optimizing CI/CD pipelines and integration tests, reducing failed deployments by 40%.
 - Developed accounting batch jobs and reporting features for large-scale financial reconciliation.
- LINE (Yahoo Group)** Tokyo, Japan
Backend Engineer *Jan 2022 – May 2023*
 - Developed high-traffic REST APIs for LINE Bank gateway systems.
 - Improved authorization and release processes, increasing test coverage from 40% to 90%.
 - Built mock servers to reduce external API dependency costs by 30%.
- Money Forward** Tokyo, Japan
Backend Engineer *Jul 2020 – Jan 2022*
 - Built financial aggregation systems integrating with 50+ banks and other financial institutions.
 - Migrated accounting services from Ruby on Rails to Spring Boot, improving throughput by 2x.
- JILIGUALA** Shanghai, China
Software Engineer *Jul 2019 – Apr 2020*
 - Developed backend systems and tooling for cross-platform educational games serving millions of users.

SELECTED PROJECTS

- claude-memory (personal)** *Feb 2026 – Present*
 Four-layer persistent-memory MCP server for Claude Code (Go, SQLite FTS5; ranking by BM25, confidence and recency; deduplication; layer promotion; 8 tools), in daily use and holding 800+ memories.
- night-shift (personal)** *Aug 2026 – Present*
 Autonomous research agent that runs scheduled headless Claude Code sessions to adversarially verify claims in AI papers, keeps only verified findings in its memory, and compiles a daily digest.

SKILLS

- Programming Languages:** Go, TypeScript, Python, Java, Kotlin
- AI Engineering:** LLM Integration (Gemini, Claude), Prompt Engineering, LLM Benchmarking, Agent Context and Memory, MCP Servers, Claude Code (Skills, Subagents, Hooks), Multi-Agent Workflows, RAG
- Backend:** Distributed Systems, APIs (gRPC, GraphQL, REST), Event-Driven Architecture (Pub/Sub, Cloud Tasks), Idempotency and Rate Limiting, Caching (Redis, Valkey), Authentication, System Design
- Cloud & Infra:** GCP (Cloud SQL, Memorystore, BigQuery, Secret Manager), Kubernetes, Istio, Terraform, Datadog Monitors as Code, GitHub Actions CI/CD, Serverless, Observability, Platform Migration
- Languages:** English (TOEIC 940), Japanese (JLPT N1), Chinese (Native)